

# Whitepaper #02: The Disconnected Pipeline

---

**Subtitle:** *Solving Dependency Management & Containerization in Secure Facilities* **Author:** Dustin J. Ober, PMP

---

## 1. Executive Summary

---

**The Bottom Line Up Front (BLUF):** The single greatest bottleneck in sovereign AI development is not compute power, but **dependency management**. Modern software engineering relies on a "connected supply chain" (PyPI, npm, Docker Hub, Hugging Face). When this chain is broken by an air-gap or secure boundary, development velocity often drops to near zero.

**The Solution:** This brief outlines a standardized "**Disconnected DevOps**" pipeline. It moves away from ad-hoc file transfers ("burning a DVD with zip files") to a structured, automated system of **Internal Mirrors** and **Containerized Artifacts**.

**The Outcome:** By implementing local package repositories (e.g., Sonatype Nexus) and utilizing container-based transport (Docker/APPTainer), organizations can restore 90% of the "connected" developer experience while maintaining 100% Zero Trust compliance.

---

## 2. The Operational Challenge: Developing in the Dark

---

In the commercial world, a "Hello World" Python script takes 30 seconds to set up. In a secure facility, it can take 3 weeks.

### 2.1 The "Dependency Hell" Cascade

A developer requests the `pandas` library. Security approves it. The developer brings it in, only to find it requires `numpy`. They bring that in, only to find it requires `pytz`.

- **The Reality:** A standard AI environment (PyTorch + LangChain + Transformers) has over **300 transitive dependencies**.
- **The Failure Mode:** Manually moving these files one by one is impossible. It leads to "Shadow IT," where developers smuggle unauthorized libraries just to get their code to run.

## 2.2 The "It Works on My Machine" Crisis

Without a shared internet connection, every workstation drifts. Developer A has `v1.2` of a library; Developer B has `v1.3`.

- **Result:** Code written on one machine crashes on another. The lack of a centralized "Truth" (like PyPI) breaks collaboration.
- 

## 3. The "Sneakernet" Protocol: Moving Assets Securely

The "Sneakernet"—physically moving data on portable media—is often ridiculed, but in closed systems, it is the only bandwidth available. The goal is to make it **robust**, not just a workaround.

### 3.1 The "Open Internet" Staging Area

You must establish a dedicated "Ingest Workstation" connected to the Open Internet. This machine does **not** contain sensitive internal data. Its only purpose is to fetch, build, and bundle public assets.

- **The "Bundle" Strategy:** Never move raw code. Move **Artifacts**.
- *Bad:* Moving a folder of `.py` scripts.
- *Good:* Moving a Docker Image (`.tar`) or a Python Wheel (`.whl`).

### 3.2 The Docker "Save/Load" Workflow

Containerization is the ultimate transport wrapper. It freezes the OS, the drivers, and the libraries into a single file.

#### Step 1: Build on the Open Internet

```
# On the Internet-connected machine
docker build -t my-ai-app:v1 .
```

**Step 2: Export to Artifact** Docker provides a native command to flatten an image into a tarball.

```
docker save -o my-ai-app_v1.tar my-ai-app:v1
```

**Step 3: Transfer & Scan** Move the `.tar` file to the transfer medium (CD/DVD/Diode). The security team scans *only this one file*. This is significantly faster than scanning thousands of loose source files.

#### Step 4: Hydrate in the Closed System

```
# On the Closed System machine
docker load -i my-ai-app_v1.tar
```

*Result:* The exact environment is restored, bit-for-bit. No missing libraries, no version conflicts.

---

## 4. Mirroring the World: The Local Repository Strategy

---

While "Sneakernet" works for moving massive artifacts (like Docker images), it is inefficient for granular dependency management. You cannot burn a DVD every time a developer needs a tiny helper library like `tqdm` or `requests`.

The solution is to establish a **Local Mirror**—an internal server that mimics the directory structure of public repositories like PyPI and Hugging Face.

### 4.1 The Infrastructure: Nexus or Artifactory

You need a "Binary Repository Manager." The two industry standards are **Sonatype Nexus** and **JFrog Artifactory**.

**The Architecture:**

1. **Open Internet Mirror:** A Nexus instance connected to the internet that proxies [pypi.org](https://pypi.org). It caches everything verified developers download.
2. **Export Mechanism:** A scheduled script exports the storage blob of the Open Internet Mirror.
3. **Closed System Mirror:** A disconnected Nexus instance inside the secure boundary where the storage blob is imported.

## 4.2 Configuring the Client (The Developer Experience)

The goal is **Transparency**. The developer should type `pip install` and have it work, without knowing the internet is gone. We achieve this by modifying the global configuration.

**For Python (pip):** Create or modify `~/.pip/pip.conf` (Linux) or `%APPDATA%\pip\pip.ini` (Windows):

```
[global]
index-url = https://nexus.internal.lab/repository/pypi-hosted/simple
trusted-host = nexus.internal.lab
```

*Result:* `pip` now ignores [pypi.org](https://pypi.org) and talks strictly to your internal server.

**For Hugging Face (Local):** Hugging Face libraries (`transformers`) attempt to ping the internet by default. You must point them to your local cache or offline endpoint using environment variables in your `.bashrc`:

```
export HF_DATASETS_OFFLINE=1
export TRANSFORMERS_OFFLINE=1
export HF_HOME=/mnt/shared/models/huggingface
```

## 5. Advanced Containerization: Apptainer (Singularity)

While Docker is the standard for *building* containers, it is often forbidden in Closed Systems (especially HPC clusters) because the Docker daemon requires `root` privileges. If a container breaks out, the attacker gains root access to the host node.

**The Solution: Apptainer** (formerly Singularity).

## 5.1 The SIF Standard

Apptainer compresses an entire container into a Single Image File (`.sif`). Unlike Docker layers, a `.sif` file is a single, immutable artifact that can be cryptographically signed.

## 5.2 The "Docker-to-Apptainer" Pipeline

You do not need to rewrite your Dockerfiles. You simply convert them at the boundary.

### Step 1: Save Docker Image (Open Internet)

```
docker save -o my-model.tar my-model:latest
```

**Step 2: Build SIF (Transfer Boundary)** Use Apptainer to convert the Docker tarball into a secure SIF image.

```
apptainer build my-model.sif docker-archive://my-model.tar
```

**Step 3: Execute Securely (Closed System)** Run the image as a standard user (no root required).

```
apptainer run --nv my-model.sif
```

*Note:* The `--nv` flag passes the NVIDIA GPU drivers from the host into the container automatically—a massive quality-of-life feature for AI workloads.

---

## 6. Security & Governance: The "Golden Image"

---

In a Closed System, you cannot "patch" vulnerabilities easily. Therefore, security shifts left—it must happen *before* the artifact enters the air-gap.

### 6.1 The Software Bill of Materials (SBOM)

Every container entering the Closed System must accompany an SBOM. This is a manifest listing every library (OS-level and Python-level) inside the image.

- **Tooling:** Use **Syft** or **Grype** to generate SBOMs during the build process.

```
syft my-model:latest -o cyclonedx-json > sbom.json
```

- **The Audit:** If a new CVE is discovered in `OpenSSL`, you query your central SBOM database to find exactly which offline containers are affected, rather than scanning terabytes of closed-system drives.

## 6.2 The "Golden Image" Strategy

Do not allow developers to bring in raw base images (like `ubuntu:latest`).

1. **Create a Base:** Security creates a "Hardened AI Base" (Ubuntu + CUDA + Python + Certs).
2. **Publish:** This image is available on the Closed System Nexus.
3. **Mandate:** All developer Dockerfiles must start with `FROM nexus.internal.lab/hardened-ai-base:v1`.

---

## 7. Conclusion

---

The "Disconnected Pipeline" is the circulatory system of a Sovereign AI capability. Without it, the hardware discussed in Whitepaper #01 is merely expensive metal.

By moving from ad-hoc file transfers to a structured architecture of **Internal Mirrors**, **Containerized Artifacts**, and **Apptainer Runtimes**, organizations can achieve a development velocity that rivals the commercial sector while adhering to the strictest security mandates. The result is a system that is secure by design, auditable by default, and resilient against supply chain attacks.

### Next in this Series:

- **Whitepaper #03:** *Private Knowledge Retrieval: Architecting Local RAG Systems.*

---

## About the Author

**Dustin J. Ober, PMP, M.Ed.** *AI Developer & Technical Instructional Systems Designer*

Dustin J. Ober is a specialist in the intersection of Artificial Intelligence, Instructional Strategy, and secure systems architecture. With a background spanning over two decades in the United States Air Force and defense contracting, he focuses on deploying high-impact technical solutions within mission-critical environments.

Unlike traditional developers who focus solely on code, Dustin bridges the gap between **technical capability** and **operational reality**. His expertise lies in architecting "Sovereign AI" systems—designing offline, air-gapped inference pipelines that allow organizations to leverage state-of-the-art intelligence without compromising data security or compliance.

He holds a Master of Education in Instructional Design & Technology and is a certified Project Management Professional (PMP). He actively develops open-source tools for the AI community, focusing on DSPy implementation, neuro-symbolic logic, and verifiable agentic workflows.

### Connect:

- **Web:** [aiober.com](http://aiober.com)
- **LinkedIn:** [linkedin.com/in/dustinober](https://linkedin.com/in/dustinober)
- **Email:** [dustinober@me.com](mailto:dustinober@me.com)

**Suggested Citation:** Ober, D. J. (2025). *The Disconnected Pipeline: Solving Dependency Management in Secure Facilities* (Whitepaper No. 02). AIOber Technical Insights.